

In [1]:

```
import tensorflow as tf
from tensorflow.keras import regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow import keras
from keras import layers
from keras import models
from tensorflow.keras import optimizers
from sklearn.metrics import classification_report
import numpy as np
```

2025-06-10 22:14:26.459062: E external/local_xla/xla/stream_executor/cuda/cuda_ff t.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
 WARNING: All log messages before absl::InitializeLog() is called are written to S TDERR
 E0000 00:00:1749590066.472617 3052959 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already bee n registered
 E0000 00:00:1749590066.476775 3052959 cuda_blas.cc:1407] Unable to register cuBLA S factory: Attempting to register factory for plugin cuBLAS when one has already been registered
 W0000 00:00:1749590066.487264 3052959 computation_placer.cc:177] computation plac er already registered. Please check linkage and avoid linking the same target mor e than once.
 W0000 00:00:1749590066.487282 3052959 computation_placer.cc:177] computation plac er already registered. Please check linkage and avoid linking the same target mor e than once.
 W0000 00:00:1749590066.487283 3052959 computation_placer.cc:177] computation plac er already registered. Please check linkage and avoid linking the same target mor e than once.
 W0000 00:00:1749590066.487284 3052959 computation_placer.cc:177] computation plac er already registered. Please check linkage and avoid linking the same target mor e than once.
 2025-06-10 22:14:26.490784: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
 To enable the following instructions: AVX2 FMA, in other operations, rebuild Tens orFlow with the appropriate compiler flags.

In [2]:

```
tf.__version__
```

Out[2]:

'2.19.0'

MODELO S

1 - Data Preprocessing

Caminhos dos sets

In [3]:

```
train_dir = 'dataset_balanceado_final/train'
validation_dir = 'dataset_balanceado_final/validation'
test_dir = 'dataset_balanceado_final/test'
```

Definir batch_size e image_size

```
In [4]: from tensorflow.keras.utils import image_dataset_from_directory

IMG_SIZE = 150
BATCH_SIZE = 32
```

Training set - É o conjunto de dados usado para treinar a rede

```
In [5]: train_dataset = image_dataset_from_directory(
    train_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical' # categorical porque temos várias classes, senão se
)
```

Found 4276 files belonging to 7 classes.

```
I0000 00:00:1749590068.836593 3052959 gpu_device.cc:2019] Created device /job:loc
alhost/replica:0/task:0/device:GPU:0 with 4804 MB memory: -> device: 0, name: NV
IDIA GeForce GTX 1660 SUPER, pci bus id: 0000:03:00.0, compute capability: 7.5
I0000 00:00:1749590068.840103 3052959 gpu_device.cc:2019] Created device /job:loc
alhost/replica:0/task:0/device:GPU:1 with 4804 MB memory: -> device: 1, name: NV
IDIA GeForce GTX 1660 SUPER, pci bus id: 0000:05:00.0, compute capability: 7.5
```

Validation set - Usado para 'testar' o modelo durante o processo de procura da melhor combinação de hiperparâmetros.

```
In [6]: validation_dataset = image_dataset_from_directory(
    validation_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical'
)
```

Found 1420 files belonging to 7 classes.

Test set - Usado para testar o modelo depois do processo de treino

```
In [7]: test_dataset = image_dataset_from_directory(
    test_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical'
)
```

Found 1420 files belonging to 7 classes.

Data Augmentation

Gerar mais dados de treino a partir de amostras de treino existentes, aumentando as amostras através de uma série de transformações aleatórias que produzem imagens com aspecto credível.

```
In [ ]: data_augmentation = keras.Sequential([
    layers.RandomFlip("horizontal"), # Efeito
    layers.RandomRotation(0.1), # rotação de 1
    layers.RandomZoom(0.2),]) # rotação de 20%
```

Métricas para avaliar os modelos

```
In [9]: # Utiliza uma função(do scikit-Learn) para avaliar o desempenho do modelo, indicando:
# f1-score do modelo
# accuracy do modelo
# accuracy por classe

from sklearn.metrics import classification_report
import numpy as np

def print_classification_metrics(model, dataset, phase_name):
    y_true = []
    y_pred = []

    for images, labels in dataset:
        preds = model.predict(images)
        y_true.extend(np.argmax(labels.numpy(), axis=1))
        y_pred.extend(np.argmax(preds, axis=1))

    print(f"\n {phase_name}")
    print(classification_report(y_true, y_pred, digits=4))
```

2 - Construir a CNN (convolucional Neural Network)

```
In [10]: # Define a camada de entrada do modelo com o formato das imagens (altura, Largura)
inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))

In [11]: #Aplica as transformações aleatórias às imagens (definidas no inicio do notebook)
x = data_augmentation(inputs)

In [12]: # Normaliza os valores dos pixels das imagens de entrada para o intervalo [0, 1]
x = layers.Rescaling(1./255)(x)
```

Passo 1 : Camada Convolucional

```
In [13]: # Primeira camada convolucional com 64 filtros 3x3 e ativação ReLU
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu", padding="same")(
# adicionado o padding="same" para garantir que o output da camada convolucional
```

Passo 2 : Camada de Pooling

```
In [14]: # Primeira camada de pooling máximo (2x2) para reduzir a dimensionalidade.
x = layers.MaxPooling2D(pool_size=2)(x)
```

Passo 3 : Adicionar mais camadas

```
In [15]: # Segunda camada convolucional com 128 filtros 3x3 e ativação ReLU.
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu", padding="same")
# Segunda camada de pooling máximo (2x2).
x = layers.MaxPooling2D(pool_size=2)(x)
```

```
In [16]: # Terceira camada convolucional com 128 filtros 3x3 e ativação ReLU.
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu", padding="same")
# Terceira camada de pooling máximo (2x2).
x = layers.MaxPooling2D(pool_size=2)(x)
```

Passo 4 : Flattening

```
In [17]: # Acha o output das camadas convolucionais para um vetor 1D.
x = layers.Flatten()(x)
```

BATCH NORMALIZATION (Facultativo)

```
In [18]: # tirar comentario desta linha abaixo se queremos usar batch normalization
# porem foi testado varias vezes e em diferentes camadas da rede, mas não melhor
#x = Layers.BatchNormalization( axis=-1, momentum=0.99, epsilon=0.001, center=True )(x)
```

Dropout (facultativo) - função de regularização

```
In [19]: # Aplica Dropout (50%) para desativar aleatoriamente neurónios, prevenindo o overfitting
x = layers.Dropout(0.5)(x)
```

Passo 5 : Full Connection

```
In [20]: # Definir função de Regularização L2 (opcional)
reg = regularizers.l2(0.01) # Executar para ativar

# Camada densa (totalmente conectada) com 128 neurónios, ativação ReLU e função de regularização
x = layers.Dense(128, activation="relu", kernel_regularizer=reg)(x)
```

Passo 6 : Output Layer

```
In [21]: # Camada de saída densa com 7 neurónios e ativação Softmax (para classificação categórica)
outputs = layers.Dense(7, activation="softmax")(x)
```

```
In [22]: # Cria o modelo Keras usando as camadas de entrada e saída definidas.
model = keras.Model(inputs=inputs, outputs=outputs)
```

3 - Treinar a rede CNN

Funções de otimização disponíveis: Adam, RMSprop e SGD

Funções de loss disponíveis: categorical_crossentropy , KLDivergence e MSE

```
In [23]: # Configura o otimizador Adam  
# Define a função de Loss: KLDivergence  
# Indica que a 'accuracy' (precisão) será a métrica durante o treino.  
model.compile(  
    optimizer=tf.keras.optimizers.Adam(),  
    loss=tf.keras.losses.KLDivergence(),  
    metrics=['accuracy'])
```

```
In [24]: #model.compile(  
#optimizer='SGD',  
#loss='mse',  
#metrics=['accuracy'])
```

```
In [25]: #model.compile(  
#optimizer=tf.keras.optimizers.RMSprop(),  
#Loss=tf.keras.Losses.KLDivergence(),  
#metrics=['accuracy'])
```

Treinar o modelo

```
In [26]: history = model.fit(  
    train_dataset, #Inicia o treino do modelo usando o conjunto de dados de trei  
    epochs=25,      # O modelo será treinado por 25 épocas (passagens completas p  
    validation_data=validation_dataset) # Usa o conjunto de dados de validação p
```

```
Epoch 1/25
```

```
I0000 00:00:1749590072.986723 3053018 cuda_dnn.cc:529] Loaded cuDNN version 90300
```

```
134/134 ━━━━━━━━━━ 14s 80ms/step - accuracy: 0.1804 - loss: 2.9115 - va  
l_accuracy: 0.2507 - val_loss: 1.9783  
Epoch 2/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.3034 - loss: 1.9019 - va  
l_accuracy: 0.3254 - val_loss: 1.8228  
Epoch 3/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.3868 - loss: 1.7169 - va  
l_accuracy: 0.4789 - val_loss: 1.5583  
Epoch 4/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.4316 - loss: 1.5989 - va  
l_accuracy: 0.5035 - val_loss: 1.4842  
Epoch 5/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.4726 - loss: 1.4959 - va  
l_accuracy: 0.4951 - val_loss: 1.4498  
Epoch 6/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.4765 - loss: 1.4581 - va  
l_accuracy: 0.4831 - val_loss: 1.5025  
Epoch 7/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.4846 - loss: 1.4575 - va  
l_accuracy: 0.4789 - val_loss: 1.5565  
Epoch 8/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.4989 - loss: 1.4091 - va  
l_accuracy: 0.5394 - val_loss: 1.3455  
Epoch 9/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5071 - loss: 1.4035 - va  
l_accuracy: 0.4887 - val_loss: 1.4702  
Epoch 10/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.4975 - loss: 1.4033 - va  
l_accuracy: 0.5458 - val_loss: 1.3391  
Epoch 11/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5175 - loss: 1.3452 - va  
l_accuracy: 0.5141 - val_loss: 1.4211  
Epoch 12/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5057 - loss: 1.3954 - va  
l_accuracy: 0.5070 - val_loss: 1.4471  
Epoch 13/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5250 - loss: 1.3703 - va  
l_accuracy: 0.5479 - val_loss: 1.3006  
Epoch 14/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5237 - loss: 1.3504 - va  
l_accuracy: 0.5099 - val_loss: 1.4027  
Epoch 15/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5286 - loss: 1.3746 - va  
l_accuracy: 0.5549 - val_loss: 1.3289  
Epoch 16/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5208 - loss: 1.3788 - va  
l_accuracy: 0.5493 - val_loss: 1.2961  
Epoch 17/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5367 - loss: 1.3235 - va  
l_accuracy: 0.5465 - val_loss: 1.3486  
Epoch 18/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5250 - loss: 1.3372 - va  
l_accuracy: 0.5387 - val_loss: 1.3425  
Epoch 19/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5391 - loss: 1.3248 - va  
l_accuracy: 0.5563 - val_loss: 1.3054  
Epoch 20/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5539 - loss: 1.2846 - va  
l_accuracy: 0.5479 - val_loss: 1.3432  
Epoch 21/25
```

```
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5411 - loss: 1.3123 - va  
l_accuracy: 0.5648 - val_loss: 1.3016  
Epoch 22/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5298 - loss: 1.3281 - va  
l_accuracy: 0.5514 - val_loss: 1.3078  
Epoch 23/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5483 - loss: 1.3064 - va  
l_accuracy: 0.5592 - val_loss: 1.3390  
Epoch 24/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5290 - loss: 1.3327 - va  
l_accuracy: 0.5535 - val_loss: 1.3448  
Epoch 25/25  
134/134 ━━━━━━━━━━ 10s 74ms/step - accuracy: 0.5280 - loss: 1.3284 - va  
l_accuracy: 0.5697 - val_loss: 1.2698
```

4 - Testar o modelo

```
In [27]: print_classification_metrics(model, test_dataset, "Modelo 1 : CNN de raiz")
```

1/1	0s	103ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	55ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	55ms/step
1/1	0s	55ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	55ms/step
1/1	0s	53ms/step
1/1	0s	53ms/step
1/1	0s	55ms/step
1/1	0s	53ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	56ms/step
1/1	0s	55ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	53ms/step
1/1	0s	55ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	82ms/step

Modelo 1 : CNN de raiz

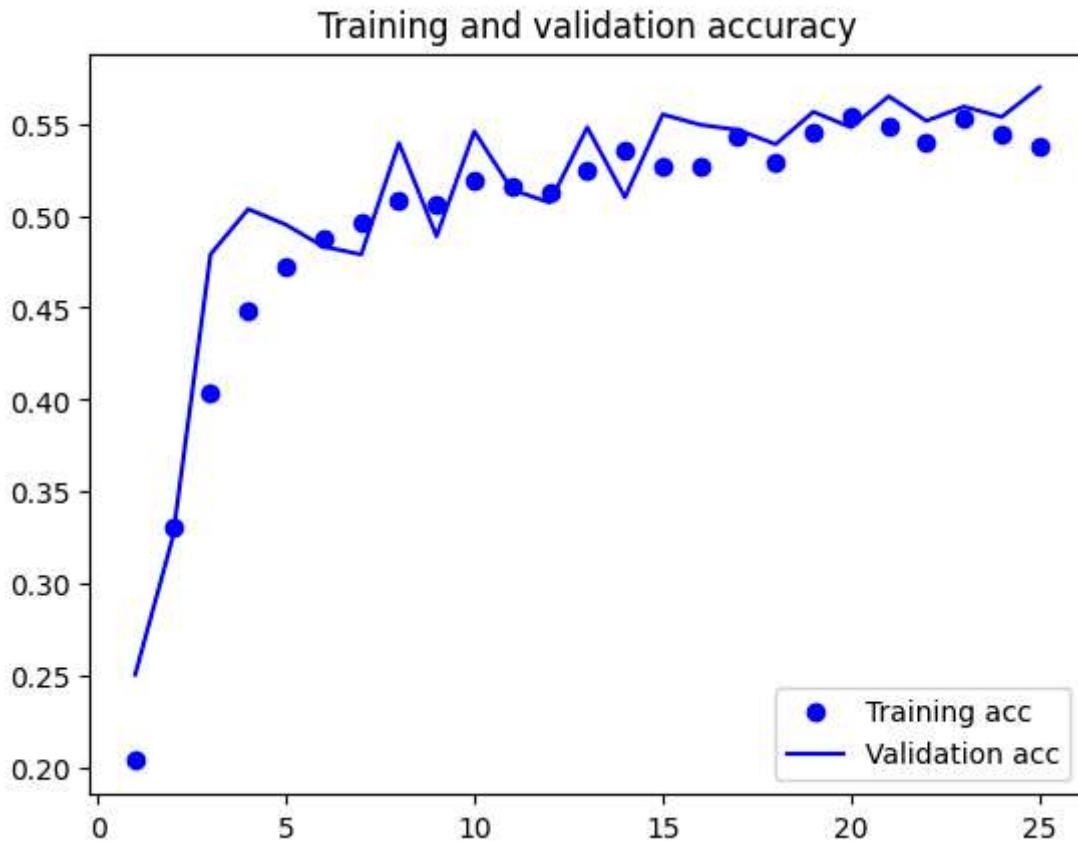
	precision	recall	f1-score	support
0	0.4602	0.4154	0.4367	195
1	0.4839	0.4545	0.4688	198
2	0.5299	0.6484	0.5832	219
3	0.5172	0.4688	0.4918	192
4	0.5311	0.5000	0.5151	222
5	0.6738	0.6300	0.6512	200
6	0.8455	0.9588	0.8986	194
accuracy			0.5817	1420
macro avg	0.5774	0.5823	0.5779	1420

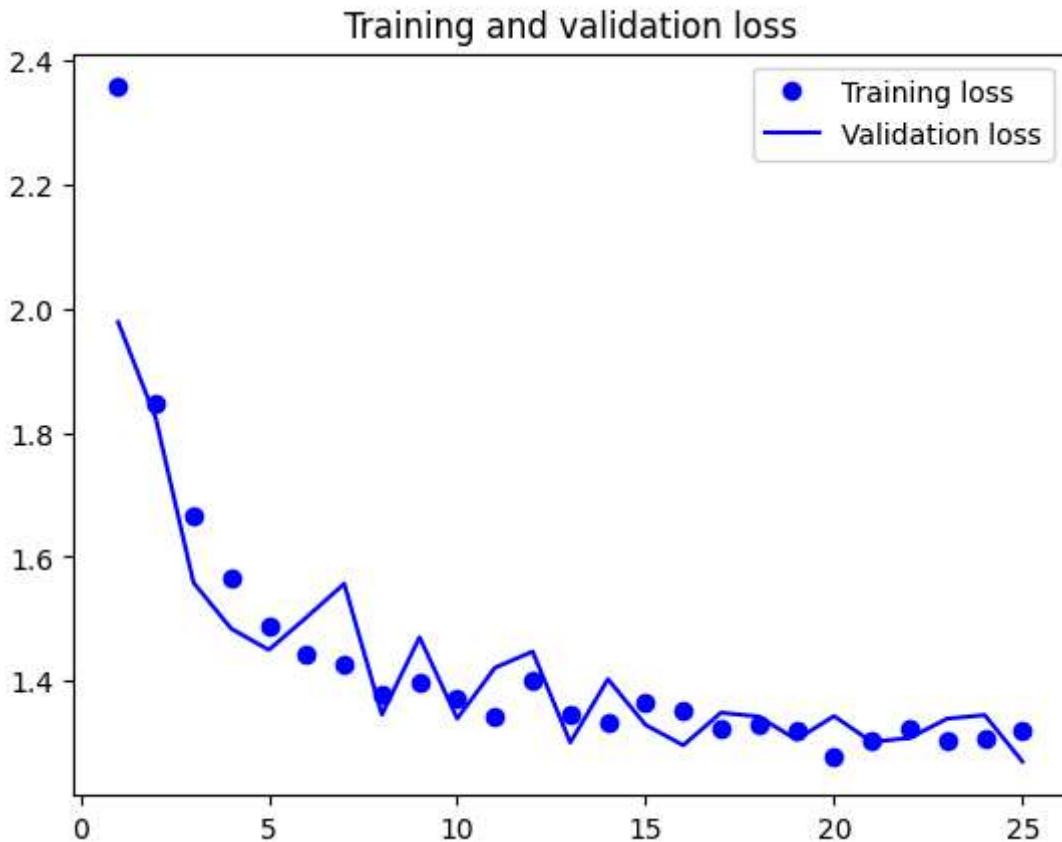
weighted avg	0.5758	0.5817	0.5768	1420
--------------	--------	--------	--------	------

2025-06-10 22:18:47.909411: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

Ver as Curvas de Loss e de Accuracy

```
In [28]: import matplotlib.pyplot as plt
accuracy = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```





Análise geral dos gráficos acima: O modelo treinou de forma "equilibrada", com desempenho consistente em treino e validação. A estabilidade das curvas sugere que o número de épocas foi adequado e o modelo está bem ajustado.

```
In [ ]: import numpy as np
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Obter previsões no test_dataset
y_true = []
y_pred = []

for images, labels in test_dataset:
    preds = model.predict(images)
    y_true.extend(np.argmax(labels.numpy(), axis=1))
    y_pred.extend(np.argmax(preds, axis=1))

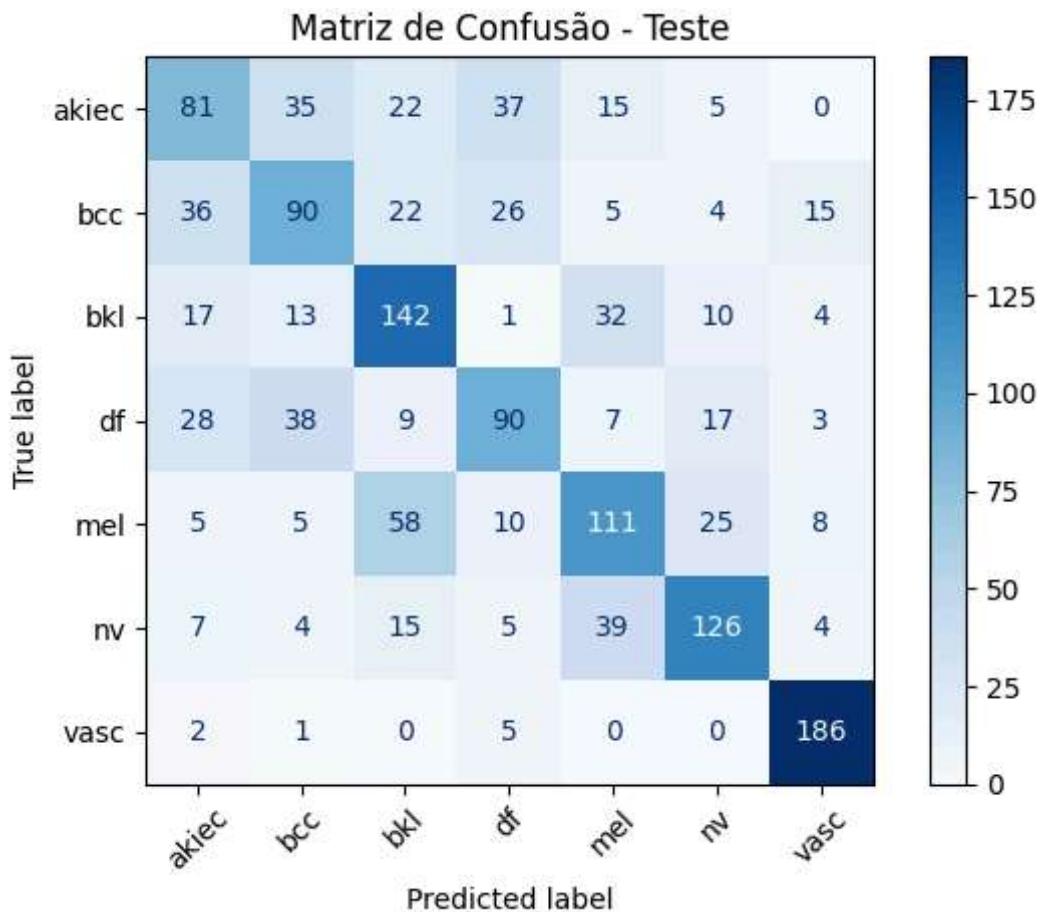
class_names = test_dataset.class_names

# Criar e mostrar a matriz de confusão
cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)

plt.figure(figsize=(10, 8))
disp.plot(cmap=plt.cm.Blues, xticks_rotation=45, values_format='d')
plt.title("Matriz de Confusão - Teste")
plt.tight_layout()
plt.show()
```

1/1	0s	89ms/step
1/1	0s	55ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	55ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	53ms/step
1/1	0s	35ms/step

```
2025-06-10 22:48:54.899190: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence <Figure size 1000x800 with 0 Axes>
```



The Kernel crashed while executing code in the current cell or a previous cell.

Please review the code in the cell(s) to identify a possible cause of the failure.

Click [here](https://aka.ms/vscodeJupyterKernelCrash) for more info.

View Jupyter [log](command:jupyter.viewOutput) for further details.

A matriz de confusão mostra melhor desempenho geral, com classes "vasc", "bkl" e "nv" muito bem classificadas. Algumas confusões persistem entre lesões visualmente semelhantes, como "akiec" e "bcc"

Salvar o modelo

```
In [29]: model.save("modeloS_2B1_com_data_aug_adam_KLD.keras")
```