

In [1]:

```
import tensorflow as tf
from tensorflow.keras import regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow import keras
from keras import layers
from keras import models
from tensorflow.keras import optimizers
from sklearn.metrics import classification_report
import numpy as np
```

2025-06-12 02:14:34.953150: E external/local\_xla/xla/stream\_executor/cuda/cuda\_ff  
t.cc:467] Unable to register cuFFT factory: Attempting to register factory for pl  
ugin cuFFT when one has already been registered  
WARNING: All log messages before absl::InitializeLog() is called are written to S  
TDERR  
E0000 00:00:1749690874.966510 91027 cuda\_dnn.cc:8579] Unable to register cuDNN  
factory: Attempting to register factory for plugin cuDNN when one has already bee  
n registered  
E0000 00:00:1749690874.970667 91027 cuda\_blas.cc:1407] Unable to register cuBLA  
S factory: Attempting to register factory for plugin cuBLAS when one has already  
been registered  
W0000 00:00:1749690874.980640 91027 computation\_placer.cc:177] computation plac  
er already registered. Please check linkage and avoid linking the same target mor  
e than once.  
W0000 00:00:1749690874.980654 91027 computation\_placer.cc:177] computation plac  
er already registered. Please check linkage and avoid linking the same target mor  
e than once.  
W0000 00:00:1749690874.980656 91027 computation\_placer.cc:177] computation plac  
er already registered. Please check linkage and avoid linking the same target mor  
e than once.  
W0000 00:00:1749690874.980657 91027 computation\_placer.cc:177] computation plac  
er already registered. Please check linkage and avoid linking the same target mor  
e than once.  
2025-06-12 02:14:34.984114: I tensorflow/core/platform/cpu\_feature\_guard.cc:210]  
This TensorFlow binary is optimized to use available CPU instructions in perfor  
mance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild Tens  
orFlow with the appropriate compiler flags.

In [2]:

```
tf.__version__
```

Out[2]:

'2.19.0'

# MODELO S

## 1 - Data Preprocessing

Caminhos dos sets

In [3]:

```
train_dir = 'dataset_balanceado_final/train'
validation_dir = 'dataset_balanceado_final/validation'
test_dir = 'dataset_balanceado_final/test'
```

Definir batch\_size e image\_size

```
In [4]: from tensorflow.keras.utils import image_dataset_from_directory

IMG_SIZE = 150
BATCH_SIZE = 32
```

Training set - É o conjunto de dados usado para treinar a rede

```
In [5]: train_dataset = image_dataset_from_directory(
    train_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical' # categorical porque temos várias classes, senão se
)
```

Found 4276 files belonging to 7 classes.

```
I0000 00:00:1749690877.404282 91027 gpu_device.cc:2019] Created device /job:loc
alhost/replica:0/task:0/device:GPU:0 with 4804 MB memory: -> device: 0, name: NV
IDIA GeForce GTX 1660 SUPER, pci bus id: 0000:03:00.0, compute capability: 7.5
I0000 00:00:1749690877.407857 91027 gpu_device.cc:2019] Created device /job:loc
alhost/replica:0/task:0/device:GPU:1 with 4804 MB memory: -> device: 1, name: NV
IDIA GeForce GTX 1660 SUPER, pci bus id: 0000:05:00.0, compute capability: 7.5
```

Validation set - Usado para 'testar' o modelo durante o processo de procura da melhor combinação de hiperparâmetros.

```
In [6]: validation_dataset = image_dataset_from_directory(
    validation_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical'
)
```

Found 1420 files belonging to 7 classes.

Test set - Usado para testar o modelo depois do processo de treino

```
In [7]: test_dataset = image_dataset_from_directory(
    test_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical'
)
```

Found 1420 files belonging to 7 classes.

## Data Augmentation

Gerar mais dados de treino a partir de amostras de treino existentes, aumentando as amostras através de uma série de transformações aleatórias que produzem imagens com aspecto credível.

```
In [8]: data_augmentation = keras.Sequential([layers.RandomFlip("horizontal"), # Efeito
                                             layers.RandomFlip("vertical"),# Aplica inv
                                             layers.RandomTranslation(0.1, 0.2),
                                             layers.RandomRotation(0.4),])
```

## Métricas para avaliar os modelos

```
In [9]: # Utiliza uma função(do scikit-Learn) para avaliar o desempenho do modelo, indi
      # f1-score do modelo
      # accuracy do modelo
      # accuracy por classe

from sklearn.metrics import classification_report
import numpy as np

def print_classification_metrics(model, dataset, phase_name):
    y_true = []
    y_pred = []

    for images, labels in dataset:
        preds = model.predict(images)
        y_true.extend(np.argmax(labels.numpy(), axis=1))
        y_pred.extend(np.argmax(preds, axis=1))

    print(f"\n {phase_name}")
    print(classification_report(y_true, y_pred, digits=4))
```

## 2 - Construir a CNN (convolucional Neural Network)

```
In [10]: # Define a camada de entrada do modelo com o formato das imagens (altura, Largur
inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))

In [11]: #Aplica as transformações aleatórias às imagens (definidas no inicio do notebook
x = data_augmentation(inputs)

In [12]: # Normaliza os valores dos pixels das imagens de entrada para o intervalo [0, 1]
x = layers.Rescaling(1./255)(x)
```

### Passo 1 : Camada Convolucional

```
In [13]: # Primeira camada convolucional com 64 filtros 3x3 e ativação ReLU
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu", padding="same")(
# adicionado o padding="same" para garantir que o output da camada convolucional
```

### Passo 2 : Camada de Pooling

```
In [14]: # Primeira camada de pooling máximo (2x2) para reduzir a dimensionalidade.
x = layers.MaxPooling2D(pool_size=2)(x)
```

### Passo 3 : Adicionar mais camadas

```
In [15]: # Segunda camada convolucional com 128 filtros 3x3 e ativação ReLU.
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu", padding="same")
# Segunda camada de pooling máximo (2x2).
x = layers.MaxPooling2D(pool_size=2)(x)
```

```
In [16]: # Terceira camada convolucional com 128 filtros 3x3 e ativação ReLU.
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu", padding="same")
# Terceira camada de pooling máximo (2x2).
x = layers.MaxPooling2D(pool_size=2)(x)
```

## Passo 4 : Flattening

```
In [17]: # Achata o output das camadas convolucionais para um vetor 1D.
x = layers.Flatten()(x)
```

BATCH NORMALIZATION (Facultativo)

```
In [18]: # tirar comentario desta linha abaixo se queremos usar batch normalization
# porem foi testado varias vezes e em diferentes camadas da rede, mas não melhor
#x = Layers.BatchNormalization( axis=-1, momentum=0.99, epsilon=0.001, center=True )(x)
```

## Dropout (facultativo) - função de regularização

```
In [19]: # Aplica Dropout (50%) para desativar aleatoriamente neurónios, prevenindo o overfitting
# Neste caso, para este modelo em específico, atingiu a pior acc sem dropout
#x = Layers.Dropout(0.5)(x)
```

## Passo 5 : Full Connection

```
In [20]: # Definir função de Regularização L2 (opcional)
reg = regularizers.l2(0.01) # Executar para ativar

# Camada densa (totalmente conectada) com 128 neurónios, ativação ReLU e função de regularização
x = layers.Dense(128, activation="relu", kernel_regularizer=reg)(x)
```

## Passo 6 : Output Layer

```
In [21]: # Camada de saída densa com 7 neurónios e ativação Softmax (para classificação categórica)
outputs = layers.Dense(7, activation="softmax")(x)
```

```
In [22]: # Cria o modelo Keras usando as camadas de entrada e saída definidas.
model = keras.Model(inputs=inputs, outputs=outputs)
```

## 3 - Treinar a rede CNN

Funções de otimização disponíveis: Adam, RMSprop e SGD

Funções de loss disponíveis: categorical\_crossentropy , KLDivergence e MSE

```
In [23]: # Configura o otimizador: SGD
# Define a função de Loss: MSE
# Indica que a 'accuracy' (precisão) será a métrica durante o treino.
model.compile(
    optimizer=tf.keras.optimizers.SGD(),
    loss='mse',
    metrics=['accuracy'])
```

```
In [24]: #model.compile(
#optimizer='SGD',
#loss='mse',
#metrics=['accuracy'])
```

```
In [25]: #model.compile(
#optimizer=tf.keras.optimizers.RMSprop(),
#Loss=tf.keras.Losses.KLDivergence(),
#metrics=['accuracy'])
```

## Treinar o modelo

```
In [26]: history = model.fit(
    train_dataset, #Inicia o treino do modelo usando o conjunto de dados de treinamento
    epochs=25,      # O modelo será treinado por 25 épocas (passagens completas pelo dataset)
    validation_data=validation_dataset) # Usa o conjunto de dados de validação para monitorar o progresso
```

Epoch 1/25

```
I0000 00:00:1749690880.720096 98298 cuda_dnn.cc:529] Loaded cuDNN version 90300
```

134/134 13s 80ms/step - accuracy: 0.1599 - loss: 2.6415 - val\_accuracy: 0.1535 - val\_loss: 2.5420  
Epoch 2/25  
134/134 10s 72ms/step - accuracy: 0.1738 - loss: 2.5099 - val\_accuracy: 0.1514 - val\_loss: 2.4157  
Epoch 3/25  
134/134 10s 72ms/step - accuracy: 0.1718 - loss: 2.3853 - val\_accuracy: 0.1444 - val\_loss: 2.2959  
Epoch 4/25  
134/134 10s 72ms/step - accuracy: 0.1667 - loss: 2.2672 - val\_accuracy: 0.1577 - val\_loss: 2.1824  
Epoch 5/25  
134/134 10s 72ms/step - accuracy: 0.1798 - loss: 2.1552 - val\_accuracy: 0.1662 - val\_loss: 2.0749  
Epoch 6/25  
134/134 10s 72ms/step - accuracy: 0.1946 - loss: 2.0490 - val\_accuracy: 0.1768 - val\_loss: 1.9729  
Epoch 7/25  
134/134 10s 72ms/step - accuracy: 0.1957 - loss: 1.9484 - val\_accuracy: 0.1993 - val\_loss: 1.8763  
Epoch 8/25  
134/134 10s 72ms/step - accuracy: 0.2141 - loss: 1.8531 - val\_accuracy: 0.2162 - val\_loss: 1.7847  
Epoch 9/25  
134/134 10s 72ms/step - accuracy: 0.2201 - loss: 1.7627 - val\_accuracy: 0.2261 - val\_loss: 1.6979  
Epoch 10/25  
134/134 10s 72ms/step - accuracy: 0.2303 - loss: 1.6770 - val\_accuracy: 0.2507 - val\_loss: 1.6156  
Epoch 11/25  
134/134 10s 72ms/step - accuracy: 0.2292 - loss: 1.5958 - val\_accuracy: 0.2465 - val\_loss: 1.5376  
Epoch 12/25  
134/134 10s 72ms/step - accuracy: 0.2309 - loss: 1.5189 - val\_accuracy: 0.2465 - val\_loss: 1.4637  
Epoch 13/25  
134/134 10s 72ms/step - accuracy: 0.2427 - loss: 1.4460 - val\_accuracy: 0.2451 - val\_loss: 1.3936  
Epoch 14/25  
134/134 10s 72ms/step - accuracy: 0.2397 - loss: 1.3768 - val\_accuracy: 0.2648 - val\_loss: 1.3272  
Epoch 15/25  
134/134 10s 72ms/step - accuracy: 0.2407 - loss: 1.3113 - val\_accuracy: 0.2732 - val\_loss: 1.2643  
Epoch 16/25  
134/134 10s 72ms/step - accuracy: 0.2563 - loss: 1.2491 - val\_accuracy: 0.2810 - val\_loss: 1.2046  
Epoch 17/25  
134/134 10s 72ms/step - accuracy: 0.2584 - loss: 1.1902 - val\_accuracy: 0.2789 - val\_loss: 1.1480  
Epoch 18/25  
134/134 10s 72ms/step - accuracy: 0.2663 - loss: 1.1344 - val\_accuracy: 0.2845 - val\_loss: 1.0944  
Epoch 19/25  
134/134 10s 72ms/step - accuracy: 0.2671 - loss: 1.0815 - val\_accuracy: 0.2908 - val\_loss: 1.0436  
Epoch 20/25  
134/134 10s 72ms/step - accuracy: 0.2781 - loss: 1.0313 - val\_accuracy: 0.2908 - val\_loss: 0.9954  
Epoch 21/25

```
134/134 ━━━━━━━━━━ 10s 72ms/step - accuracy: 0.2796 - loss: 0.9838 - va  
l_accuracy: 0.2866 - val_loss: 0.9497  
Epoch 22/25  
134/134 ━━━━━━━━━━ 10s 72ms/step - accuracy: 0.2691 - loss: 0.9387 - va  
l_accuracy: 0.2859 - val_loss: 0.9064  
Epoch 23/25  
134/134 ━━━━━━━━━━ 10s 72ms/step - accuracy: 0.2804 - loss: 0.8960 - va  
l_accuracy: 0.2901 - val_loss: 0.8654  
Epoch 24/25  
134/134 ━━━━━━━━━━ 10s 72ms/step - accuracy: 0.2656 - loss: 0.8555 - va  
l_accuracy: 0.2873 - val_loss: 0.8265  
Epoch 25/25  
134/134 ━━━━━━━━━━ 10s 72ms/step - accuracy: 0.2782 - loss: 0.8170 - va  
l_accuracy: 0.2859 - val_loss: 0.7896
```

## 4 - Testar o modelo

```
In [27]: print_classification_metrics(model, test_dataset, "Modelo 1 : CNN de raiz")
```

1/1	0s	103ms/step
1/1	0s	56ms/step
1/1	0s	55ms/step
1/1	0s	53ms/step
1/1	0s	57ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	54ms/step
1/1	0s	56ms/step
1/1	0s	56ms/step
1/1	0s	53ms/step
1/1	0s	58ms/step
1/1	0s	54ms/step
1/1	0s	55ms/step
1/1	0s	57ms/step
1/1	0s	61ms/step
1/1	0s	59ms/step
1/1	0s	55ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	57ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	58ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	57ms/step
1/1	0s	54ms/step
1/1	0s	55ms/step
1/1	0s	55ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	53ms/step
1/1	0s	54ms/step
1/1	0s	54ms/step
1/1	0s	57ms/step
1/1	0s	54ms/step
1/1	0s	84ms/step

### Modelo 1 : CNN de raiz

	precision	recall	f1-score	support
0	0.2807	0.4923	0.3575	195
1	0.2057	0.4040	0.2726	198
2	0.5000	0.0183	0.0352	219
3	0.0000	0.0000	0.0000	192
4	0.2807	0.4640	0.3497	222
5	0.3917	0.6150	0.4786	200
6	0.0000	0.0000	0.0000	194
accuracy			0.2859	1420
macro avg	0.2370	0.2848	0.2134	1420

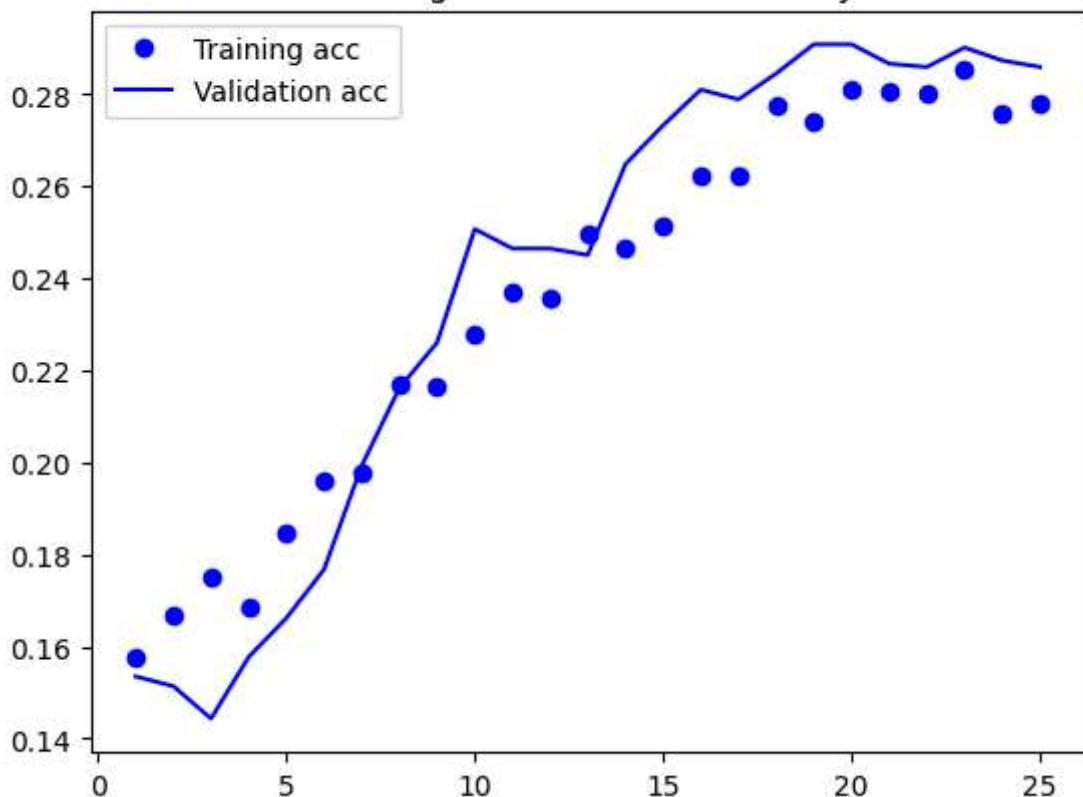
weighted avg	0.2434	0.2859	0.2146	1420
--------------	--------	--------	--------	------

```
2025-06-12 02:18:48.451744: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
/home/eliana/escola/venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/eliana/escola/venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/eliana/escola/venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

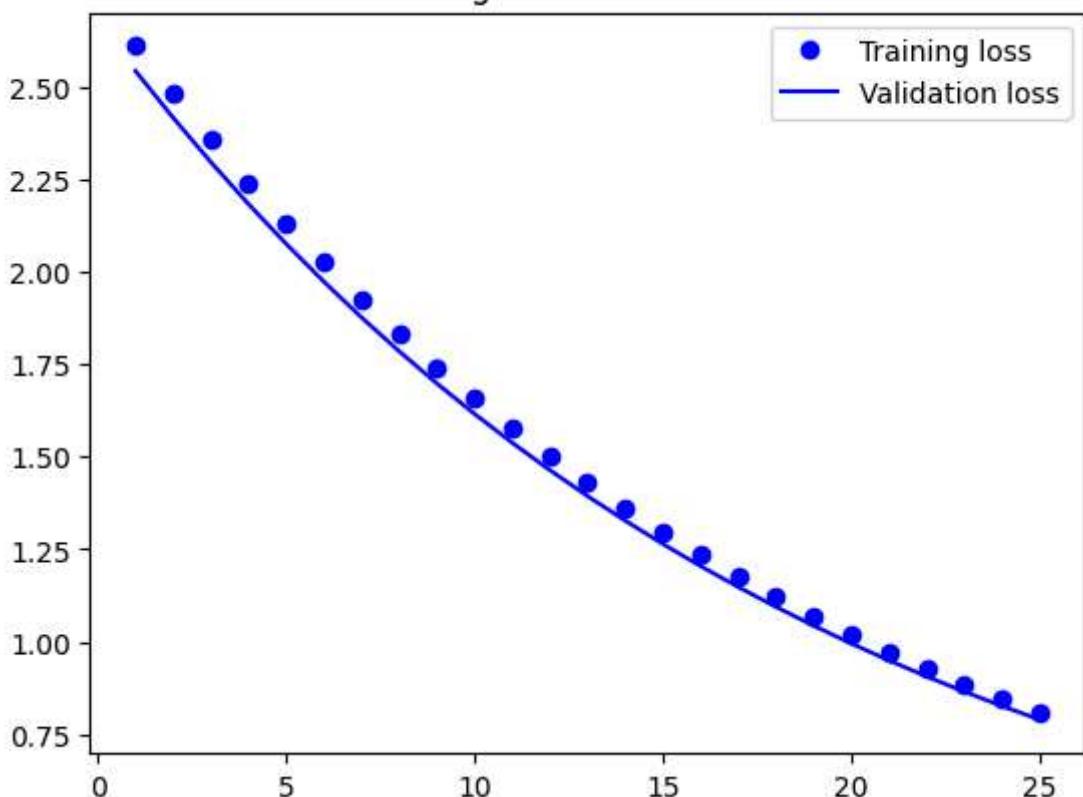
## Ver as Curvas de Loss e de Accuracy

```
In [28]: import matplotlib.pyplot as plt
accuracy = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

### Training and validation accuracy



### Training and validation loss



A acurácia de treino e validação aumentam de forma gradual ao longo das 25 épocas, atingindo cerca de 0.29, o que indica que o modelo está a aprender, mas ainda tem desempenho bastante limitado. Isso pode estar relacionado ao uso da MSE como função de perda, que não é ideal para classificação.

As perdas de treino e validação diminuem de forma consistente e estão muito próximas, o que indica que não há overfitting, mas sim que o modelo está a aprender de forma ineficiente para a tarefa proposta. A função de perda MSE pode estar a dificultar a convergência para uma solução que realmente separe bem as classes.

## Matriz de Confusão

```
In [30]: import numpy as np
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Obter previsões no test_dataset
y_true = []
y_pred = []

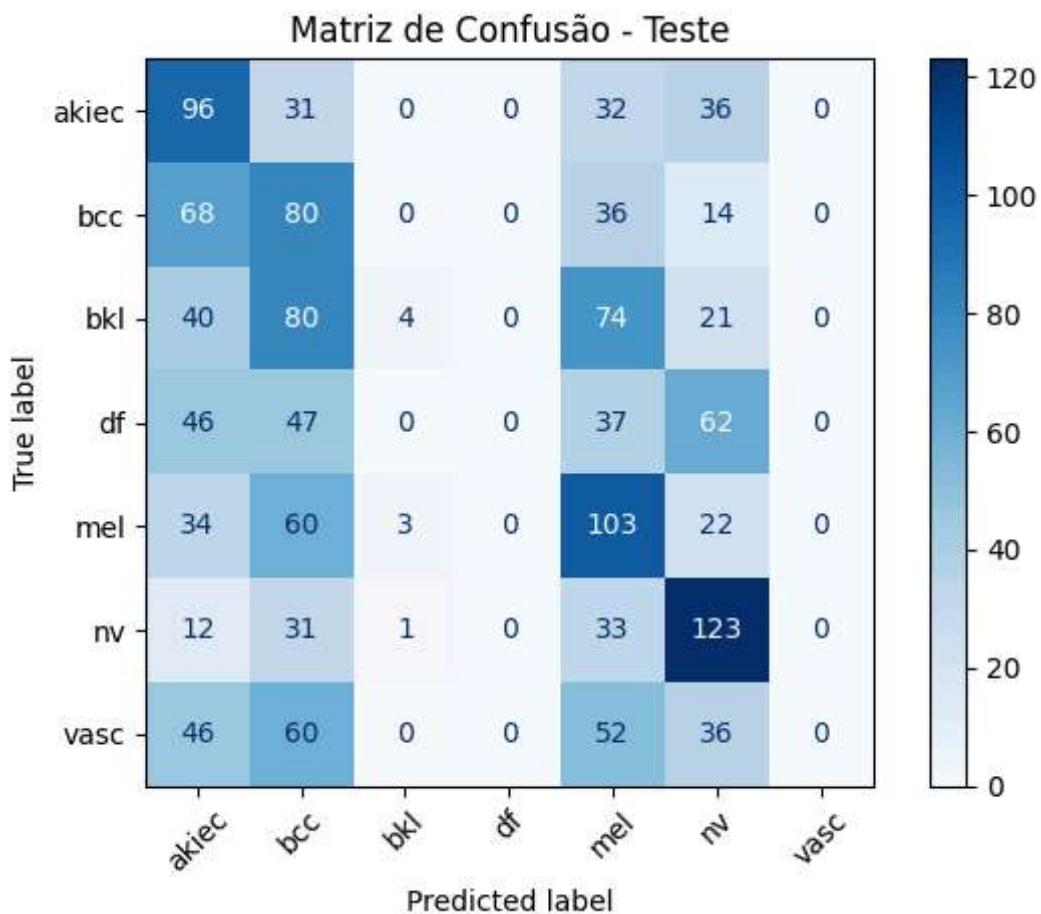
for images, labels in test_dataset:
    preds = model.predict(images)
    y_true.extend(np.argmax(labels.numpy(), axis=1))
    y_pred.extend(np.argmax(preds, axis=1))

class_names = test_dataset.class_names

# Criar e mostrar a matriz de confusão
cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)

plt.figure(figsize=(10, 8))
disp.plot(cmap=plt.cm.Blues, xticks_rotation=45, values_format='d')
plt.title("Matriz de Confusão - Teste")
plt.tight_layout()
plt.show()
```

```
2025-06-12 02:18:52.545861: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence <Figure size 1000x800 with 0 Axes>
```



## Salvar o modelo

```
In [29]: model.save("modeloS_2B3_com_data_aug_SGD_MSE_worst_acc.keras")
```