

```
In [1]: import tensorflow as tf
from tensorflow.keras import regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow import keras
from keras import layers
from keras import models
from tensorflow.keras import optimizers
from sklearn.metrics import classification_report
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
import numpy as np
```

```
2025-06-11 23:21:04.520766: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1749680464.534134 15362 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1749680464.538076 15362 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
W0000 00:00:1749680464.547928 15362 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1749680464.547944 15362 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1749680464.547945 15362 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1749680464.547946 15362 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
2025-06-11 23:21:04.551293: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
```

```
In [2]: tf.__version__
```

```
Out[2]: '2.19.0'
```

MODELO S - Optuna

O presente notebook foi testado, numa segunda abordagem, com as configurações obtidas pelo optuna

1 - Data Preprocessing

Caminhos dos sets

```
In [3]: train_dir = 'dataset_balanceado_final/train'
validation_dir = 'dataset_balanceado_final/validation'
test_dir = 'dataset_balanceado_final/test'
```

Definir batch_size e image_size

```
In [4]: from tensorflow.keras.utils import image_dataset_from_directory

IMG_SIZE = 150
BATCH_SIZE = 32
```

Training set - É o conjunto de dados usado para treinar a rede

```
In [5]: train_dataset = image_dataset_from_directory(
    train_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical' # categorical porque temos várias classes, senão se
)
```

Found 4276 files belonging to 7 classes.

```
I0000 00:00:1749680466.967290 15362 gpu_device.cc:2019] Created device /job:loc
alhost/replica:0/task:0/device:GPU:0 with 4804 MB memory: -> device: 0, name: NV
IDIA GeForce GTX 1660 SUPER, pci bus id: 0000:03:00.0, compute capability: 7.5
I0000 00:00:1749680466.970832 15362 gpu_device.cc:2019] Created device /job:loc
alhost/replica:0/task:0/device:GPU:1 with 4804 MB memory: -> device: 1, name: NV
IDIA GeForce GTX 1660 SUPER, pci bus id: 0000:05:00.0, compute capability: 7.5
```

Validation set - Usado para 'testar' o modelo durante o processo de procura da melhor combinação de hiperparâmetros.

```
In [6]: validation_dataset = image_dataset_from_directory(
    validation_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical'
)
```

Found 1420 files belonging to 7 classes.

Test set - Usado para testar o modelo depois do processo de treino

```
In [7]: test_dataset = image_dataset_from_directory(
    test_dir,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    label_mode='categorical'
)
```

Found 1420 files belonging to 7 classes.

Data Augmentation

Gerar mais dados de treino a partir de amostras de treino existentes, aumentando as amostras através de uma série de transformações aleatórias que produzem imagens com aspecto credível.

```
In [8]: data_augmentation = keras.Sequential([layers.RandomFlip("horizontal"), # Efeito
                                             layers.RandomRotation(0.1),
                                             layers.RandomTranslation(0.1, 0.2),
                                             layers.RandomFlip("vertical"), # rotação
                                             layers.RandomZoom(0.2),]) # rotação de 20%
```

Métricas para avaliar os modelos

```
In [9]: # Utiliza uma função(do scikit-learn) para avaliar o desempenho do modelo, indicando
# f1-score do modelo
# accuracy do modelo
# accuracy por classe

from sklearn.metrics import classification_report
import numpy as np

def print_classification_metrics(model, dataset, phase_name):
    y_true = []
    y_pred = []

    for images, labels in dataset:
        preds = model.predict(images)
        y_true.extend(np.argmax(labels.numpy(), axis=1))
        y_pred.extend(np.argmax(preds, axis=1))

    print(f"\n {phase_name}")
    print(classification_report(y_true, y_pred, digits=4))
```

2 - Construir a CNN (convolucional Neural Network)

```
In [10]: # Define a camada de entrada do modelo com o formato das imagens (altura, Largura e profundidade)
inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))

In [11]: #Aplica as transformações aleatórias às imagens (definidas no inicio do notebook)
x = data_augmentation(inputs)

In [12]: # Normaliza os valores dos pixels das imagens de entrada para o intervalo [0, 1]
x = layers.Rescaling(1./255)(x)
```

Passo 1 : Camada Convolucional

```
In [13]: # Primeira camada convolucional com 64 filtros 3x3 e ativação ReLU
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu", padding="same")(
    # adicionado o padding="same" para garantir que o output da camada convolucional
    # seja o mesmo que a input
```

Passo 2 : Camada de Pooling

```
In [14]: # Primeira camada de pooling máximo (2x2) para reduzir a dimensionalidade.
x = layers.MaxPooling2D(pool_size=2)(x)
```

Passo 3 : Adicionar mais camadas

```
In [15]: # Segunda camada convolucional com 128 filtros 3x3 e ativação ReLU.
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu", padding="same")
# Segunda camada de pooling máximo (2x2).
x = layers.MaxPooling2D(pool_size=2)(x)
```

```
In [16]: # Terceira camada convolucional com 128 filtros 3x3 e ativação ReLU.
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu", padding="same")
# Terceira camada de pooling máximo (2x2).
x = layers.MaxPooling2D(pool_size=2)(x)
```

Passo 4 : Flattening

```
In [17]: # Achata o output das camadas convolucionais para um vetor 1D.
x = layers.Flatten()(x)
```

BATCH NORMALIZATION (Facultativo)

```
In [18]: # tirar comentario desta linha abaixo se queremos usar batch normalization
# porem foi testado varias vezes e em diferentes camadas da rede, mas não melhor
#x = Layers.BatchNormalization( axis=-1, momentum=0.99, epsilon=0.001, center=True )(x)
```

Dropout (optuna disse para não usar) - função de regularização

```
In [19]: # Aplica Dropout (50%) para desativar aleatoriamente neurónios, prevenindo o overfitting
x = layers.Dropout(0.5)(x)
```

Passo 5 : Full Connection

```
In [20]: # Definir função de Regularização L2 dada pelo optuna
reg = regularizers.l2(0.0002426218) # Executar para ativar

# Camada densa (totalmente conectada) com 128 neurónios, ativação ReLU e função de regularização
x = layers.Dense(128, activation="relu", kernel_regularizer=reg)(x)
```

Passo 6 : Output Layer

```
In [21]: # Camada de saída densa com 7 neurónios e ativação Softmax (para classificação categórica)
outputs = layers.Dense(7, activation="softmax")(x)
```

```
In [22]: # Cria o modelo Keras usando as camadas de entrada e saída definidas.
model = keras.Model(inputs=inputs, outputs=outputs)
```

3 - Treinar a rede CNN

Funções de otimização disponíveis: Adam, RMSprop e SGD

Funções de loss disponíveis: categorical_crossentropy , KLDivergence e MSE

```
In [23]: # Configura o otimizador Adam com Learning rate de 0.001
# Define a função de loss: Categorical_Crossentropy com label smoothing de 0.095
# Indica que a 'accuracy' (precisão) será a métrica durante o treino.
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss=tf.keras.losses.CategoricalCrossentropy(label_smoothing=0.095),
    metrics=['accuracy'])
```

Treinar o modelo

```
In [ ]: #utilizar early stopping -> o modelo pára de treinar se a accuracy não subir durante o treinamento
callbacks = [
    EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True),
    ModelCheckpoint(filepath='models/melhor_modelo.h5', save_best_only=True, monitor='val_loss')
]

history = model.fit(
    train_dataset, #Inicia o treino do modelo usando o conjunto de dados de treinamento
    epochs=100, # Apesar de ter um numero alto aqui, o EarlyStopping para automaticamente quando a accuracy não subir mais
    validation_data=validation_dataset # Usa o conjunto de dados de validação para monitorar a performance
    #, callbacks=callbacks
)
```

Epoch 1/100

I0000 00:00:1749680471.338111 15429 cuda_dnn.cc:529] Loaded cuDNN version 90300

```
134/134 ━━━━━━━━━━ 15s 81ms/step - accuracy: 0.1748 - loss: 2.0181 - va  
l_accuracy: 0.2782 - val_loss: 1.7705  
Epoch 2/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.3391 - loss: 1.7492 - va  
l_accuracy: 0.4120 - val_loss: 1.5985  
Epoch 3/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.4252 - loss: 1.5830 - va  
l_accuracy: 0.4106 - val_loss: 1.5637  
Epoch 4/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.4626 - loss: 1.5161 - va  
l_accuracy: 0.4676 - val_loss: 1.4421  
Epoch 5/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5031 - loss: 1.4521 - va  
l_accuracy: 0.5035 - val_loss: 1.4243  
Epoch 6/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5061 - loss: 1.4401 - va  
l_accuracy: 0.5458 - val_loss: 1.3827  
Epoch 7/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5218 - loss: 1.3938 - va  
l_accuracy: 0.5500 - val_loss: 1.3593  
Epoch 8/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5228 - loss: 1.4072 - va  
l_accuracy: 0.4768 - val_loss: 1.4665  
Epoch 9/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5308 - loss: 1.3875 - va  
l_accuracy: 0.5634 - val_loss: 1.3369  
Epoch 10/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5318 - loss: 1.3517 - va  
l_accuracy: 0.5458 - val_loss: 1.3706  
Epoch 11/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5464 - loss: 1.3581 - va  
l_accuracy: 0.5606 - val_loss: 1.3318  
Epoch 12/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5568 - loss: 1.3248 - va  
l_accuracy: 0.5718 - val_loss: 1.3243  
Epoch 13/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5522 - loss: 1.3495 - va  
l_accuracy: 0.5739 - val_loss: 1.3238  
Epoch 14/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5581 - loss: 1.3305 - va  
l_accuracy: 0.6021 - val_loss: 1.2958  
Epoch 15/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5571 - loss: 1.3373 - va  
l_accuracy: 0.5817 - val_loss: 1.2872  
Epoch 16/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5643 - loss: 1.3186 - va  
l_accuracy: 0.5965 - val_loss: 1.2928  
Epoch 17/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5805 - loss: 1.2985 - va  
l_accuracy: 0.5979 - val_loss: 1.2779  
Epoch 18/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5807 - loss: 1.2909 - va  
l_accuracy: 0.5908 - val_loss: 1.2716  
Epoch 19/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5873 - loss: 1.2815 - va  
l_accuracy: 0.6070 - val_loss: 1.2713  
Epoch 20/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5999 - loss: 1.2567 - va  
l_accuracy: 0.5824 - val_loss: 1.2820  
Epoch 21/100
```

```
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5980 - loss: 1.2617 - va  
l_accuracy: 0.6042 - val_loss: 1.2639  
Epoch 22/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.5928 - loss: 1.2677 - va  
l_accuracy: 0.6007 - val_loss: 1.2729  
Epoch 23/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6094 - loss: 1.2511 - va  
l_accuracy: 0.6049 - val_loss: 1.2426  
Epoch 24/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6105 - loss: 1.2400 - va  
l_accuracy: 0.6092 - val_loss: 1.2515  
Epoch 25/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6257 - loss: 1.2173 - va  
l_accuracy: 0.6254 - val_loss: 1.2185  
Epoch 26/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6153 - loss: 1.2322 - va  
l_accuracy: 0.6275 - val_loss: 1.2234  
Epoch 27/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6206 - loss: 1.2268 - va  
l_accuracy: 0.6070 - val_loss: 1.2528  
Epoch 28/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6235 - loss: 1.2248 - va  
l_accuracy: 0.6359 - val_loss: 1.2228  
Epoch 29/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6240 - loss: 1.2158 - va  
l_accuracy: 0.5993 - val_loss: 1.3083  
Epoch 30/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6425 - loss: 1.2015 - va  
l_accuracy: 0.6521 - val_loss: 1.1834  
Epoch 31/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6386 - loss: 1.2059 - va  
l_accuracy: 0.5979 - val_loss: 1.2627  
Epoch 32/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6338 - loss: 1.2047 - va  
l_accuracy: 0.6465 - val_loss: 1.2058  
Epoch 33/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6170 - loss: 1.2219 - va  
l_accuracy: 0.6451 - val_loss: 1.2148  
Epoch 34/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6321 - loss: 1.2100 - va  
l_accuracy: 0.6141 - val_loss: 1.2400  
Epoch 35/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6495 - loss: 1.1978 - va  
l_accuracy: 0.6394 - val_loss: 1.2274  
Epoch 36/100  
134/134 ━━━━━━━━━━ 10s 75ms/step - accuracy: 0.6363 - loss: 1.1933 - va  
l_accuracy: 0.6437 - val_loss: 1.1966  
Epoch 37/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6475 - loss: 1.1891 - va  
l_accuracy: 0.5965 - val_loss: 1.2918  
Epoch 38/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6354 - loss: 1.2030 - va  
l_accuracy: 0.6549 - val_loss: 1.1831  
Epoch 39/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6594 - loss: 1.1837 - va  
l_accuracy: 0.6394 - val_loss: 1.2312  
Epoch 40/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6556 - loss: 1.1970 - va  
l_accuracy: 0.6648 - val_loss: 1.1855  
Epoch 41/100
```

```
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6590 - loss: 1.1661 - va  
l_accuracy: 0.6599 - val_loss: 1.2110  
Epoch 42/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6434 - loss: 1.1821 - va  
l_accuracy: 0.6775 - val_loss: 1.1814  
Epoch 43/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6592 - loss: 1.1757 - va  
l_accuracy: 0.6838 - val_loss: 1.1587  
Epoch 44/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6682 - loss: 1.1512 - va  
l_accuracy: 0.6655 - val_loss: 1.2150  
Epoch 45/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6679 - loss: 1.1551 - va  
l_accuracy: 0.6500 - val_loss: 1.2279  
Epoch 46/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6791 - loss: 1.1419 - va  
l_accuracy: 0.6739 - val_loss: 1.1616  
Epoch 47/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6695 - loss: 1.1627 - va  
l_accuracy: 0.6887 - val_loss: 1.1709  
Epoch 48/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6564 - loss: 1.1748 - va  
l_accuracy: 0.6725 - val_loss: 1.1905  
Epoch 49/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6796 - loss: 1.1693 - va  
l_accuracy: 0.6683 - val_loss: 1.2044  
Epoch 50/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6880 - loss: 1.1456 - va  
l_accuracy: 0.6521 - val_loss: 1.2240  
Epoch 51/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6753 - loss: 1.1631 - va  
l_accuracy: 0.6930 - val_loss: 1.1706  
Epoch 52/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6605 - loss: 1.1686 - va  
l_accuracy: 0.6908 - val_loss: 1.1512  
Epoch 53/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6930 - loss: 1.1304 - va  
l_accuracy: 0.6754 - val_loss: 1.1759  
Epoch 54/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6862 - loss: 1.1500 - va  
l_accuracy: 0.6718 - val_loss: 1.1785  
Epoch 55/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6954 - loss: 1.1204 - va  
l_accuracy: 0.6887 - val_loss: 1.1691  
Epoch 56/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6879 - loss: 1.1361 - va  
l_accuracy: 0.6746 - val_loss: 1.1915  
Epoch 57/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6845 - loss: 1.1407 - va  
l_accuracy: 0.6796 - val_loss: 1.1783  
Epoch 58/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7066 - loss: 1.1148 - va  
l_accuracy: 0.6930 - val_loss: 1.1503  
Epoch 59/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7081 - loss: 1.1269 - va  
l_accuracy: 0.6993 - val_loss: 1.1605  
Epoch 60/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6917 - loss: 1.1193 - va  
l_accuracy: 0.6831 - val_loss: 1.1800  
Epoch 61/100
```

```
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6855 - loss: 1.1436 - va  
l_accuracy: 0.6887 - val_loss: 1.1632  
Epoch 62/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7031 - loss: 1.1100 - va  
l_accuracy: 0.6930 - val_loss: 1.1473  
Epoch 63/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7150 - loss: 1.1021 - va  
l_accuracy: 0.6662 - val_loss: 1.2047  
Epoch 64/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7134 - loss: 1.1109 - va  
l_accuracy: 0.6944 - val_loss: 1.1584  
Epoch 65/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7119 - loss: 1.1157 - va  
l_accuracy: 0.6782 - val_loss: 1.1821  
Epoch 66/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7087 - loss: 1.1146 - va  
l_accuracy: 0.6901 - val_loss: 1.1649  
Epoch 67/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7119 - loss: 1.1010 - va  
l_accuracy: 0.6915 - val_loss: 1.1663  
Epoch 68/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7026 - loss: 1.1073 - va  
l_accuracy: 0.7092 - val_loss: 1.1285  
Epoch 69/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7089 - loss: 1.1034 - va  
l_accuracy: 0.6704 - val_loss: 1.2017  
Epoch 70/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7046 - loss: 1.1256 - va  
l_accuracy: 0.7141 - val_loss: 1.1351  
Epoch 71/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7121 - loss: 1.1014 - va  
l_accuracy: 0.7148 - val_loss: 1.1266  
Epoch 72/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.6988 - loss: 1.1108 - va  
l_accuracy: 0.7092 - val_loss: 1.1421  
Epoch 73/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7202 - loss: 1.1019 - va  
l_accuracy: 0.6894 - val_loss: 1.1971  
Epoch 74/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7127 - loss: 1.1073 - va  
l_accuracy: 0.6923 - val_loss: 1.1692  
Epoch 75/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7108 - loss: 1.1019 - va  
l_accuracy: 0.7085 - val_loss: 1.1196  
Epoch 76/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7009 - loss: 1.0935 - va  
l_accuracy: 0.6915 - val_loss: 1.1916  
Epoch 77/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7150 - loss: 1.0960 - va  
l_accuracy: 0.7113 - val_loss: 1.1245  
Epoch 78/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7117 - loss: 1.0911 - va  
l_accuracy: 0.6880 - val_loss: 1.1849  
Epoch 79/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7095 - loss: 1.0970 - va  
l_accuracy: 0.7007 - val_loss: 1.1471  
Epoch 80/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7226 - loss: 1.0791 - va  
l_accuracy: 0.7049 - val_loss: 1.1348  
Epoch 81/100
```

```
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7356 - loss: 1.0672 - va  
l_accuracy: 0.7070 - val_loss: 1.1377  
Epoch 82/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7312 - loss: 1.0692 - va  
l_accuracy: 0.6704 - val_loss: 1.1812  
Epoch 83/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7281 - loss: 1.0852 - va  
l_accuracy: 0.6894 - val_loss: 1.1922  
Epoch 84/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7324 - loss: 1.0706 - va  
l_accuracy: 0.6845 - val_loss: 1.1902  
Epoch 85/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7136 - loss: 1.0872 - va  
l_accuracy: 0.7014 - val_loss: 1.1579  
Epoch 86/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7259 - loss: 1.0738 - va  
l_accuracy: 0.7183 - val_loss: 1.1384  
Epoch 87/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7328 - loss: 1.0762 - va  
l_accuracy: 0.7070 - val_loss: 1.1374  
Epoch 88/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7316 - loss: 1.0678 - va  
l_accuracy: 0.7021 - val_loss: 1.1637  
Epoch 89/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7241 - loss: 1.0811 - va  
l_accuracy: 0.7148 - val_loss: 1.1379  
Epoch 90/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7211 - loss: 1.0976 - va  
l_accuracy: 0.6852 - val_loss: 1.1960  
Epoch 91/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7287 - loss: 1.0831 - va  
l_accuracy: 0.7021 - val_loss: 1.1529  
Epoch 92/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7253 - loss: 1.0875 - va  
l_accuracy: 0.6845 - val_loss: 1.1971  
Epoch 93/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7259 - loss: 1.0765 - va  
l_accuracy: 0.7113 - val_loss: 1.1364  
Epoch 94/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7251 - loss: 1.0769 - va  
l_accuracy: 0.7063 - val_loss: 1.1537  
Epoch 95/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7413 - loss: 1.0630 - va  
l_accuracy: 0.6930 - val_loss: 1.1774  
Epoch 96/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7440 - loss: 1.0644 - va  
l_accuracy: 0.7134 - val_loss: 1.1474  
Epoch 97/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7258 - loss: 1.0715 - va  
l_accuracy: 0.7014 - val_loss: 1.1705  
Epoch 98/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7459 - loss: 1.0593 - va  
l_accuracy: 0.7000 - val_loss: 1.1623  
Epoch 99/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7518 - loss: 1.0428 - va  
l_accuracy: 0.7014 - val_loss: 1.1733  
Epoch 100/100  
134/134 ━━━━━━━━━━ 10s 76ms/step - accuracy: 0.7338 - loss: 1.0658 - va  
l_accuracy: 0.7021 - val_loss: 1.1522
```

4 - Testar o modelo

```
In [25]: print_classification_metrics(model, test_dataset, "Modelo 1 : CNN de raiz")
```

```

1/1 ━━━━━━ 0s 104ms/step
1/1 ━━━━━━ 0s 54ms/step
1/1 ━━━━━━ 0s 55ms/step
1/1 ━━━━━━ 0s 55ms/step
1/1 ━━━━━━ 0s 54ms/step
1/1 ━━━━━━ 0s 55ms/step
1/1 ━━━━━━ 0s 55ms/step
1/1 ━━━━━━ 0s 54ms/step
1/1 ━━━━━━ 0s 56ms/step
1/1 ━━━━━━ 0s 54ms/step
1/1 ━━━━━━ 0s 56ms/step
1/1 ━━━━━━ 0s 54ms/step
1/1 ━━━━━━ 0s 85ms/step

```

Modelo 1 : CNN de raiz

	precision	recall	f1-score	support
0	0.5694	0.6308	0.5985	195
1	0.6923	0.5455	0.6102	198
2	0.6172	0.5890	0.6028	219
3	0.7336	0.8750	0.7981	192
4	0.5826	0.6351	0.6078	222
5	0.7733	0.6650	0.7151	200
6	0.9592	0.9691	0.9641	194
accuracy			0.6972	1420
macro avg	0.7040	0.7014	0.6995	1420

weighted avg	0.7002	0.6972	0.6956	1420
--------------	--------	--------	--------	------

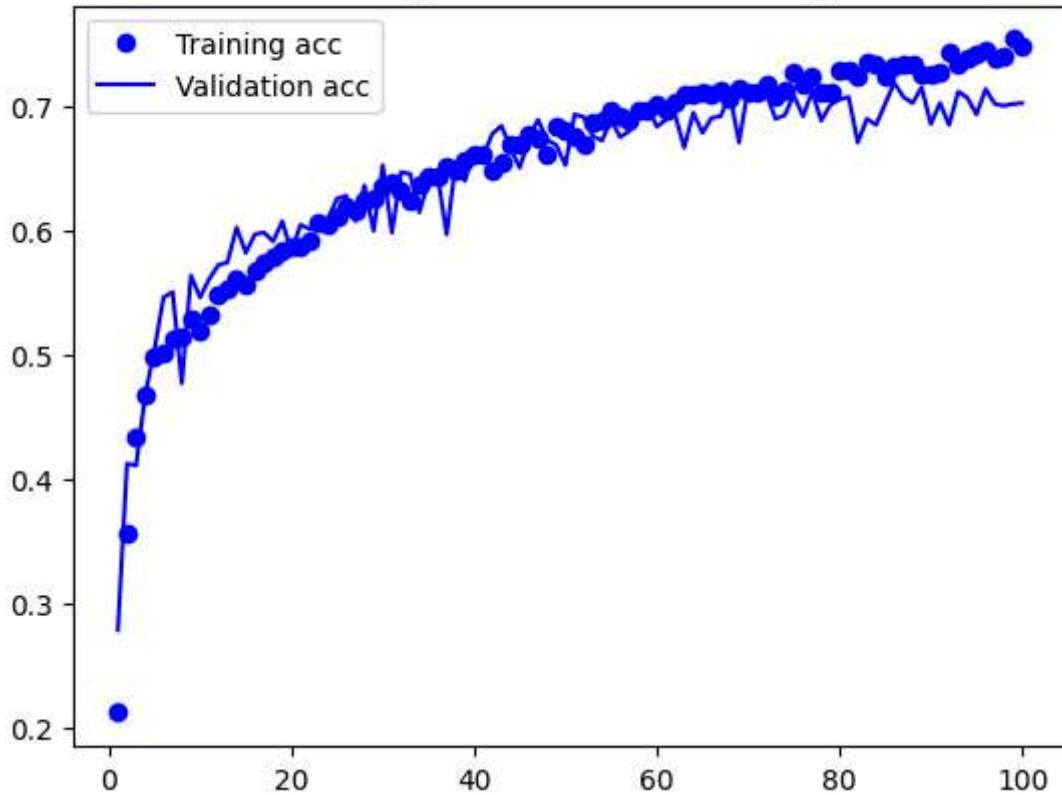
2025-06-11 23:38:11.586496: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

Curvas de Loss e de Accuracy

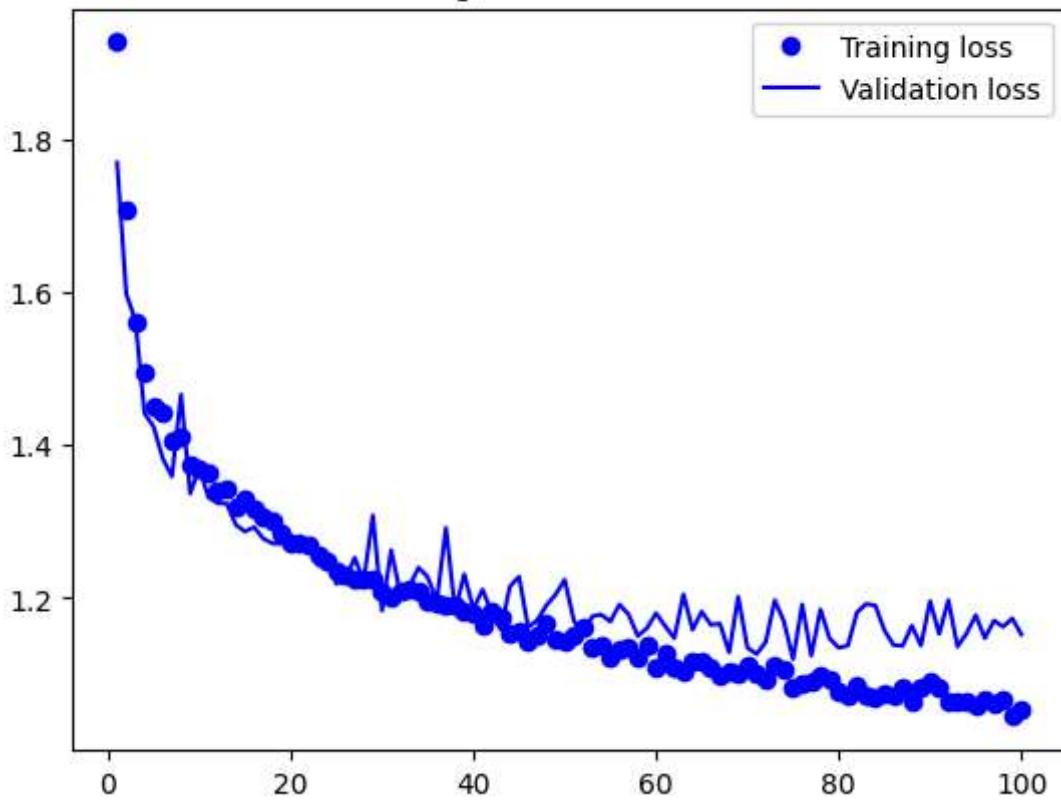
In [26]:

```
import matplotlib.pyplot as plt
accuracy = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

Training and validation accuracy



Training and validation loss



A acurácia de treino e validação aumentam de forma consistente ao longo das épocas, com valores próximos e sem grande divergência, o que indica um bom ajuste do modelo sem sinais claros de overfitting.

As curvas de loss de treino e validação diminuem gradualmente, mantendo-se próximas e estáveis ao longo do tempo. Isso mostra que o modelo está a aprender de forma eficiente e generaliza bem para os dados de validação.

Matriz de Confusão

```
In [27]: import numpy as np
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Obter previsões no test_dataset
y_true = []
y_pred = []

for images, labels in test_dataset:
    preds = model.predict(images)
    y_true.extend(np.argmax(labels.numpy(), axis=1))
    y_pred.extend(np.argmax(preds, axis=1))

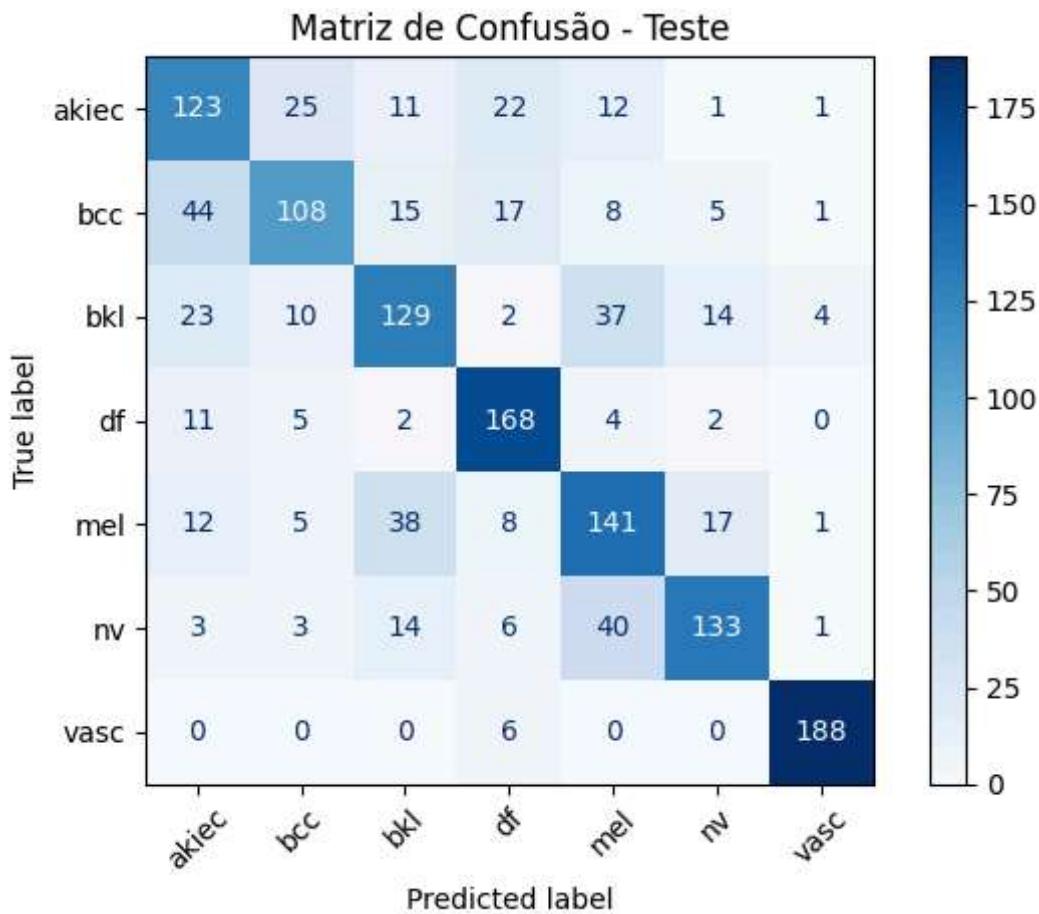
class_names = test_dataset.class_names

# Criar e mostrar a matriz de confusão
cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)

plt.figure(figsize=(10, 8))
```

```
disp.plot(cmap=plt.cm.Blues, xticks_rotation=45, values_format='d')
plt.title("Matriz de Confusão - Teste")
plt.tight_layout()
plt.show()
```

```
2025-06-11 23:38:15.473451: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence <Figure size 1000x800 with 0 Axes>
```



O modelo apresenta um bom desempenho global, com elevada taxa de acerto nas classes 'df', 'nv' e 'vasc'. No entanto, observa-se confusão significativa entre 'akiec', 'bcc' e 'bkl', o que pode dever-se à semelhança visual entre estas classes.

Salvar o modelo

```
In [ ]: model.save("modelS_2C_optuna_com_data_aug_adam_cat_cross.keras")
```

The Kernel crashed while executing code in the current cell or a previous cell.
Please review the code in the cell(s) to identify a possible cause of the failure.
Click [here](https://aka.ms/vscodeJupyterKernelCrash) for more info.
View Jupyter [log](command:jupyter.viewOutput) for further details.